HX-Citadel Ansible Infrastructure Analysis

Analyzed by: DeepAgent **Date:** October 12, 2025 **Repository:** hx-citadel-ansible

Analysis Scope: Core Ansible infrastructure, playbooks, roles, inventory, and configuration

Table of Contents

- 1. Executive Summary
- 2. Project Statistics
- 3. Architecture Overview
- 4. Configuration Management
- 5. Inventory Structure
- 6. Variable Hierarchy
- 7. Playbook Organization
- 8. Role Architecture
- 9. Scripts and Utilities
- 10. Strengths and Best Practices
- 11. Areas for Improvement
- 12. Critical Recommendations
- 13. Security Considerations
- 14. Operational Excellence

Executive Summary

The HX-Citadel Ansible infrastructure is a **well-organized, production-ready deployment framework** for a distributed Al/LLM orchestration platform. The repository demonstrates strong adherence to Ansible best practices with a clear separation of concerns, modular role design, and comprehensive variable management.

Key Highlights

Strengths:

- Excellent FQDN-based infrastructure design preventing hardcoded IPs
- Well-structured role hierarchy with 24 specialized roles
- Comprehensive variable layering (global → group → host)
- Modular playbook design with clear deployment stages
- Strong testing framework integration
- Good documentation coverage

Areas Requiring Attention:

- Vault file management needs consolidation (vault.yml.broken indicates issues)
- Some inventory hosts use different SSH keys (inconsistency)

- Missing role-level README.md files for documentation
- No explicit rollback/recovery playbooks (except MCP recovery)
- Potential IP hardcoding violations need verification

Infrastructure Scope

- 17 managed hosts across 8 host groups
- 24 custom roles + 2 external Galaxy roles
- 26 playbooks covering deployment, validation, and recovery
- 217 YAML files across the entire project
- 9 utility scripts for automation and validation

Project Statistics

Total YAML Files: 217 Custom Roles: 24

Galaxy Roles: 2 (PostgreSQL, Redis)

Main Playbooks: 26
Inventory Files: 2
Group Variable Files: 3
Host Variable Files: 3
Utility Scripts: 9

Test Files: Multiple (see TESTS_ANALYSIS.md)

Role Breakdown

Category	Roles	Purpose
Base Infrastructure	base-setup	System configuration, log- ging, health checks
Database Layer	postgresql, redis, postgresql_role, redis_role	Data persistence and caching
Vector Database	qdrant, qdrant_web_ui	Embeddings and similarity search
LLM Services	ollama, litellm	Model hosting and API gate- way
Orchestration	orchestrator_* (9 roles)	Core application logic and workflow
API Layer	fastapi, prisma	Application interfaces
Observability	observability	Monitoring and metrics
MCP Server	fastmcp_server	Model context protocol

Architecture Overview

```
graph TB
    subgraph "Entry Point"
       A[site.yml]
    subgraph "Deployment Stages"
        B[deploy-base.yml]
        C[deploy-db.yml]
        D[deploy-vector.yml]
        E[deploy-llm.yml]
        F[deploy-api.yml]
        G[deploy-orchestrator-local.yml]
    end
    subgraph "Host Groups"
       H[db_nodes]
        I[vector_nodes]
        J[llm_nodes]
       K[orchestrator_nodes]
        L[api_nodes]
       M[ui_nodes]
       N[mcp_nodes]
    end
    subgraph "Roles Layer"
       0[base-setup]
        P[postgresql/redis]
        Q[qdrant]
        R[ollama/litellm]
        S[fastapi/prisma]
        T[orchestrator_*]
    end
   A --> B
   A --> C
   A --> D
   A --> E
   A --> F
   A --> G
   B --> 0
   C --> P
   C --> H
   D --> Q
   D --> I
   E --> R
   E --> J
   F --> S
   F --> K
   F --> L
   G --> T
    G --> K
```

Infrastructure Layers

- 1. Base Layer: Common configuration, logging, health checks
- 2. **Data Layer**: PostgreSQL (app DB) + Redis (caching)
- 3. Vector Layer: Qdrant for embeddings and similarity search

- 4. **LLM Layer**: Ollama (model hosting) + LiteLLM (unified API)
- 5. Application Layer: FastAPI, Prisma ORM
- 6. Orchestration Layer: Core business logic with multiple framework integrations
- 7. **UI/Frontend Layer**: Web interfaces
- 8. MCP Layer: Model Context Protocol servers

Configuration Management

ansible.cfg

Location: /home/ubuntu/hx-citadel-ansible/ansible.cfg

```
[defaults]
inventory = ./inventory/prod.ini
roles_path = ./roles
host_key_checking = False
retry_files_enabled = False
gathering = smart
fact_caching = jsonfile
fact_caching_connection = ./.ansible_cache
fact_caching_timeout = 86400
vault_password_file = ~/.ansible_vault_pass

[ssh_connection]
pipelining = True
ssh_args = -o ControlMaster=auto -o ControlPersist=60m -o ControlPath=~/.ssh/cm-%r@%h:
%p
```

Strengths V

- Smart gathering with fact caching (24-hour timeout) reduces execution time
- SSH multiplexing enabled for performance optimization
- Pipelining enabled for reduced SSH overhead
- No retry files prevents clutter
- Vault password file configured for automated operations

Concerns 1

- host_key_checking = False is a **security risk** in production
- Vulnerable to MITM attacks
- Should only be used in isolated/trusted networks
- Recommendation: Enable host key checking and maintain known_hosts

site.yml - Master Playbook

```
---
- import_playbook: playbooks/deploy-base.yml
- import_playbook: playbooks/deploy-db.yml
- import_playbook: playbooks/deploy-vector.yml
- import_playbook: playbooks/deploy-llm.yml
- import_playbook: playbooks/deploy-api.yml
- import_playbook: playbooks/deploy-orchestrator-local.yml
```

Design Pattern: Sequential import strategy ensures proper dependency ordering.

Evaluation:

- Clear separation of deployment stages
- V Easy to run individual stages via tags or direct playbook execution
- No error handling between stages
- No pre/post deployment validation hooks
- A Consider adding --check mode support verification

requirements.yml - External Dependencies

```
roles:
    src: geerlingguy.postgresql
    name: postgresql_role
    src: geerlingguy.redis
    name: redis_role

collections:
    name: community.docker
    name: community.general
```

Evaluation:

- **U**ses trusted community roles (geerlingguy)
- Version pinning would be beneficial
- Missing version constraints (e.g., version: "3.4.0")
- No specification of collection versions

Recommendation:

```
roles:
    - src: geerlingguy.postgresql
    version: "3.5.0" # Pin to specific version
    name: postgresql_role
```

Inventory Structure

Production Inventory (prod.ini)

Location: inventory/prod.ini

Host Groups

Group	Hosts	Purpose
all	17 hosts	Global scope
db_nodes	hx-sqldb-server	PostgreSQL + Redis
vector_nodes	hx-vectordb-server	Qdrant vector DB
llm_nodes	hx-ollama1, hx-ollama2	Model hosting
orchestrator_nodes	hx-orchestrator-server	Core orchestration
api_nodes	hx-litellm-server, hx-prisma- server	API services
monitoring_nodes	hx-metrics-server	Observability
mcp_nodes	hx-mcp1-server	MCP protocol
ui_nodes	hx-qwebui-server	Web UI

Global Variables

```
[all:vars]
ansible_user=agent0
ansible_ssh_private_key_file=~/.ssh/id_rsa
ansible_python_interpreter=/usr/bin/python3
ansible_connection=ssh
ansible_timeout=30
```

Strengths 🔽

- Clear host-to-group mappings
- Consistent naming convention (hx-*-server)
- Logical grouping by function/service
- Global SSH settings for consistency

Issues 🛕

Inconsistent SSH Key Usage:

```
# Most hosts use id_rsa
ansible_ssh_private_key_file=~/.ssh/id_rsa

# But these two use ed25519
hx-qwebui-server ansible_host=192.168.10.53 ansible_ssh_private_key_file=~/.ssh/id_ed25519
hx-mcp1-server ansible_host=192.168.10.59 ansible_ssh_private_key_file=~/.ssh/id_ed25519
```

Why this matters:

- Creates confusion about which key to use

- May indicate these hosts were added later
- Could lead to authentication failures if keys aren't synchronized

Recommendation:

- Standardize on one key type (preferably ed25519 for security)
- OR document why certain hosts require different keys
- Consider using ansible_ssh_private_key_file in group_vars instead

StrictHostKeyChecking Disabled:

```
ansible_ssh_common_args='-o StrictHostKeyChecking=no'
```

- This is a security vulnerability
- Only justified for initial provisioning or testing
- · Should be removed in production

Variable Hierarchy

Ansible evaluates variables in this precedence order (highest to lowest):

```
    Extra vars (-e on CLI)
    Task vars
    Block vars
    Role vars
    Play vars
    Host vars
    Group vars
    Facts/registered vars
    Role defaults
```

Current Structure

```
group_vars/
all/
fqdn_map.yml # FQDN mappings (EXCELLENT!)
main.yml # Global defaults
vault.yml # Encrypted secrets
vault.yml.broken # Problem indicator
bd_nodes.yml # Database-specific vars
llm_nodes.yml # LLM-specific vars
host_vars/
hx-ollama1.yml # Ollama1-specific models
hx-orchestrator-server.yml # Orchestrator-specific config
```

group vars/all/main.yml

```
ubuntu_version: "24.04"
python_version: "3.12"
app_dir: "/opt/hx-citadel-shield"
log_dir: "/var/log/hx-citadel"
services_to_restart: []
```

Evaluation:

- Clean, minimal global defaults
- Appropriate abstraction level
- - services_to_restart: [] seems unused (verify)

group_vars/all/fqdn_map.yml

This file is **OUTSTANDING** and represents a best practice for infrastructure management.

Key Features:

1. Four mapping dictionaries:

```
    hx_hosts_fqdn : Short hostname → FQDN
    hx_hosts_ip : FQDN → IP (for reporting only)
    hx_ip_to_fqdn : IP → FQDN (for auto-fix tasks)
    hx hosts short ip : Short hostname → IP (legacy)
```

2. Clear usage guidance in comments:

```
# Short hostname → FQDN mapping
# Usage: {{ hx_hosts_fqdn['hx-ollama1'] }}
```

1. Prevents hardcoded IPs/hostnames in:

- Templates
- Configuration files
- Application code
- Service definitions

Example Usage:

```
# Instead of:
postgresql_host: "192.168.10.48" # BAD

# Use:
postgresql_host: "{{ hx_hosts_fqdn['hx-sqldb-server'] }}" # GOOD
```

Why This Matters:

- Makes infrastructure changes easy (update one file, not 100+ templates)
- Enables environment promotion (dev \rightarrow staging \rightarrow prod)
- Reduces errors from typos in hostnames/IPs
- Facilitates disaster recovery and migration

group vars/db nodes.yml

```
db_name: "hx_citadel"
db_owner: "hx_pg"
db_state: "present"
db_encoding: "UTF8"
db_lc_collate: "en_US.UTF-8"
db_lc_ctype: "en_US.UTF-8"

pg_login_unix_socket: "/var/run/postgresql"
pg_login_user: "postgres"

postgresql_listen_address: "0.0.0.0"
postgresql_port: 5432

redis_bind_interface: "127.0.0.1 {{ redis_additional_bind_ip | default(ansible_default_ipv4.address) }}"
redis_additional_bind_ip: "192.168.10.48"
```

Issues:

1. Hardcoded IP:

Should be:

```
redis_additional_bind_ip: "{{ hx_hosts_ip[hx_hosts_fqdn['hx-sqldb-server']] }}"
```

1. Security Concern:

```
postgresql_listen_address: "0.0.0.0" # Listens on ALL interfaces
```

- This exposes PostgreSQL to all network interfaces
- Should restrict to specific interfaces or use firewall rules
- Consider: "127.0.0.1,{{ ansible_default_ipv4.address }}"

1. Missing Redis Password:

```
# redis_requirepass: "CHANGEME_STRONG" # Commented out!
```

- Redis has no authentication enabled
- Critical security vulnerability if exposed to network
- Must be configured via vault

host_vars Examples

Evaluation:

- ✓ Appropriate use of host_vars for host-specific model lists
- Clean, simple structure
- No model version pinning (could cause inconsistencies)

Recommendation:

```
ollama_models:
    name: gemma3
    version: "27b"
    checksum: "sha256:abc123..." # For verification
    name: gpt-oss
    version: "20b"
```

Vault Management

Issue Found:

```
group_vars/all/vault.yml
group_vars/all/vault.yml.broken # 🛕 What happened here?
```

Critical Questions:

- 1. Why does vault.yml.broken exist?
- 2. Was there a vault password issue?
- 3. Are there backup/recovery procedures for vault files?
- 4. Is vault.yml properly encrypted?

Recommendation:

- Document vault recovery procedures in README.md
- Remove .broken file or move to a backups/ directory
- Implement vault file rotation strategy
- Consider using separate vaults per environment

Playbook Organization

Overview

Location: playbooks/
Total Playbooks: 26

Categories:

- Deployment (16 playbooks)
- Validation (3 playbooks)

- Recovery (1 playbook)
- Maintenance (indirect via maintenance/ directory)

Deployment Playbooks

Playbook	Purpose	Target Hosts
deploy-base.yml	System baseline configuration	all
deploy-db.yml	PostgreSQL + Redis setup	db_nodes
deploy-vector.yml	Qdrant vector database	vector_nodes
deploy-llm.yml	Ollama model hosting	llm_nodes
deploy-api.yml	LiteLLM + Prisma API layers	api_nodes
deploy-orchestrator-loc- al.yml	Local orchestrator setup	orchestrator_nodes

Orchestrator Sub-Playbooks

The orchestrator has 8 specialized sub-playbooks:

- 1. deploy-orchestrator-base.yml Foundation setup
- 2. deploy-orchestrator-fastapi.yml FastAPI application
- 3. deploy-orchestrator-postgresql.yml Database connections
- 4. deploy-orchestrator-redis.yml Cache layer
- 5. deploy-orchestrator-qdrant.yml Vector DB integration
- 6. deploy-orchestrator-lightrag.yml LightRAG framework
- 7. deploy-orchestrator-workers.yml Background workers
- 8. deploy-orchestrator-langgraph.yml LangGraph integration

Additional Framework Integrations:

- deploy-copilotkit.yml CopilotKit integration
- deploy-pydantic-ai.yml Pydantic AI framework

Validation Playbooks

Playbook	Purpose	When to Run
preflight-check.yml	Pre-deployment validation	Before any deployment
validate-mcp-prereqs.yml	MCP prerequisites check	Before MCP deployment
smoke-tests.yml	Post-deployment health checks	After deployment
deploy-with-validation.yml	Deployment with inline validation	Full stack deployment

Recovery Playbooks

• recover-mcp-server.yml (14,858 lines!) - Comprehensive MCP recovery procedures

Example: deploy-db.yml

```
- name: Deploy database services
 hosts: db nodes
 become: yes
  roles:
    - postgresql
    - redis
  post_tasks:
    - name: Verify PostgreSQL connectivity
      community.postgresql.postgresql_ping:
        db: "{{ db_name | default('hx_citadel') }}"
        login_user: "{{ db_owner | default('hx_pg') }}"
        login_password: "{{ db_password | default(omit) }}"
        login_unix_socket: "{{ pg_login_unix_socket | default('/var/run/postgresql') }
}"
      become: yes
      become_user: postgres
      changed_when: false
      no_log: true
  tags:
    - db
```

Strengths :

- Uses post tasks for validation
- Properly uses become user for PostgreSQL operations
- no_log: true prevents password leakage
- changed when: false for idempotent checks
- Tagged for selective execution

Recommendation:

Add pre tasks for validation:

Role Architecture

Role Inventory

Total Roles: 24 custom + 2 Galaxy

Main Task Files: 76

Role Structure Example: base-setup

```
roles/base-setup/
─ defaults/
                                 # Default variable values
      └─ main.yml
    handlers/
                                  # Event handlers (restart services, etc.)
      └─ main.yml
   tasks/
     main.yml # Main task entry point

06-logrotate.yml # Log rotation config

07-sudo.yml # Sudo permissions
Ш
П
\square
      □ 08-convenience-scripts.yml # Helper scripts
09-health-check.yml # Health monitoring
    templates/
     activate-shield.sh.j2  # Environment activation script
health-check.sh.j2  # Health check script
logrotate.conf.j2  # Log rotation config
manage-services.sh.j2  # Service management
sudoers.j2  # Sudo configuration
\overline{\square}
vars/
      └─ main.yml
                                              # Role-specific variables
```

Design Patterns:

- Tasks numbered for execution order clarity (06, 07, 08, 09)
- Inja2 templates for all configuration files
- Separate handlers for service management
- Clear separation of defaults vs. vars

Role Structure Example: ollama

Observation:

- Much simpler structure than base-setup
- No templates or defaults directory
- Suggests straightforward installation without heavy configuration

Role Documentation Gap

Critical Issue: Most roles lack README.md files

Current State:

```
$ find roles -name "README.md" | wc -l
# Likely 0 or very few
```

Impact:

- New team members can't understand role purposes quickly
- No documentation of role variables
- No usage examples
- Difficult to maintain

Recommendation:

Create standardized README.md template:

```
# Role: role_name
## Purpose
Brief description of what this role does.
## Requirements
- OS: Ubuntu 24.04
- Dependencies: list here
## Role Variables
### Required Variables
- `var_name`: Description
### Optional Variables
- `var_name`: Description (default: value)
## Dependencies
List other roles that must run first.
## Example Playbook
\`\`\`yaml
- hosts: servers
 roles:
    - role: role_name
     var_name: value
/'/'/
## Tags
- `tag_name`: Description
## Author
Team name / contact
```

Scripts and Utilities

Script Inventory

Location: scripts/ **Total Scripts:** 9

Script	Purpose	Туре
check-fqdn.sh	Validates FQDN usage across codebase	Linting
auto-fix-types.sh	Automated type correction	Auto-fix
validate-types.sh	Type checking validation	Linting
fix-main-py-indentation.sh	Python indentation fixer	Auto-fix
slack-notify.sh	Slack notifications	Integration
linear-api-call.sh	Linear API interactions	Integration
fetch-linear-issue.sh	Fetch Linear issues	Integration
test-linear-issue.sh	Test Linear API connectivity	Integration
fix-linear-issue.sh	Automated issue fixing	Integration

Integration Ecosystem

The scripts reveal **tight integration** with:

- Slack Notifications and alerts
- Linear Issue tracking and automation
- **Type checking** Code quality enforcement
- FQDN validation Infrastructure consistency

check-fqdn.sh Analysis

This script is **critical** for maintaining FQDN discipline:

Purpose:

- Scans all YAML files for hardcoded IPs/hostnames
- Validates against .fqdn-allowlist
- Enforces the use of hx_hosts_fqdn variables

Recommendation:

- Integrate into pre-commit hooks
- Run in CI/CD pipeline
- Document exceptions in .fqdn-allowlist

Slack Integration

Files:

- scripts/slack-notify.sh
- Likely used for deployment notifications

Potential Issues:

- Are Slack webhooks stored in vault?
- Is there documentation on notification triggers?
- Can notifications be disabled for testing?

Linear Integration

Purpose: Automated issue management

Files:

- linear-api-call.sh Base API wrapper
- fetch-linear-issue.sh Retrieve issue details
- fix-linear-issue.sh Automated issue resolution (8,583 lines!)
- test-linear-issue.sh API connectivity verification

Questions:

- How are Linear API tokens managed? (should be in vault)
- What triggers automated issue fixes?
- Are there safeguards against accidental issue modifications?

Strengths and Best Practices

1. FQDN-Based Infrastructure ****

Outstanding implementation. The fqdn_map.yml file with its four mapping dictionaries is a **gold standard** for infrastructure management.

Benefits:

- Single source of truth for host information
- Easy environment migration
- Prevents hardcoded values
- Facilitates disaster recovery

2. Modular Role Design ******

Well-architected role separation:

- Clear single-responsibility principle
- Reusable components
- Orchestrator broken into 9 specialized sub-roles
- External Galaxy roles for common services

3. Variable Hierarchy ******

Proper use of group_vars and host_vars:

- Global defaults in group vars/all/
- Group-specific configs in group_vars/<group>.yml
- Host-specific overrides in host vars/

<host>.yml

- Clean separation of concerns

4. Playbook Organization *

Logical structure:

- Master site.yml imports all deployment stages
- Each stage can run independently
- Post-task validation in key playbooks
- Tags for selective execution

Testing Framework ★★★★★

As documented in TESTS ANALYSIS.md:

- Comprehensive unit tests
- Integration tests
- Load test planning
- Type checking integration
- 5,665 lines of test code

6. Validation Playbooks ★★★★

Proactive validation:

- Preflight checks before deployment
- Smoke tests after deployment
- MCP prerequisite validation
- Inline validation playbook option

7. SSH Optimization ****

Performance tuning:

- ControlMaster/ControlPersist for connection reuse
- Pipelining enabled
- Smart fact gathering with caching
- 24-hour fact cache timeout

8. Automation Scripts ***

Quality enforcement:

- FQDN validation automation
- Type checking automation
- Auto-fix capabilities
- Integration with external tools

9. Security Consciousness ***

Security features:

- Vault for secrets management
- no log: true in sensitive tasks
- Unix socket authentication for PostgreSQL
- Separate sudo configuration

10. Clear Naming Conventions ★★★★

Consistent patterns:

- All hosts prefixed with hx-

- Role names clearly indicate purpose
- Playbook names follow deploy-<service>.yml pattern
- Group names match service types

Areas for Improvement

1. Vault File Management 1. HIGH PRIORITY

Issue:

```
group_vars/all/vault.yml
group vars/all/vault.yml.broken # What happened?
```

Problems:

- Unclear vault status
- No documented recovery procedures
- Possible encryption/password issues
- Risk of losing encrypted secrets

Recommendations:

1. Immediate Actions:

- Document what happened with vault.yml.broken
- Verify vault.yml is properly encrypted
- Test vault password access
- Remove or move .broken file

1. Long-term Solutions:

- Implement vault backup strategy
- Document vault recovery procedures in README
- Consider separate vaults per environment
- Set up vault password rotation schedule
- Use ansible-vault rekey for password changes

2. Documentation Template:

```
## Vault Management
### Vault Password Location
- Password file: `~/.ansible vault pass`
- Backup location: [specify secure location]
### Rotating Vault Password
\`\`\`bash
ansible-vault rekey group vars/all/vault.yml
1,1,1,
### Viewing Vault Contents
\`\`\`bash
ansible-vault view group_vars/all/vault.yml
1,1,1,
### Emergency Recovery
If vault password is lost:
1. [document recovery steps]
2. [contact information]
```

2. Inconsistent SSH Key Usage 1

Issue:

```
# Most hosts
ansible_ssh_private_key_file=~/.ssh/id_rsa

# Exceptions
hx-qwebui-server ansible_ssh_private_key_file=~/.ssh/id_ed25519
hx-mcpl-server ansible_ssh_private_key_file=~/.ssh/id_ed25519
```

Problems:

- Confusion about key management
- Potential authentication failures
- Unclear why certain hosts differ

Recommendations:

1. Standardize on ed25519:

- More secure than RSA
- Better performance
- Industry best practice

1. Implementation:

```
# In group_vars/all/main.yml
ansible_ssh_private_key_file: "~/.ssh/id_ed25519"
# Remove overrides from inventory unless documented
```

1. If exceptions are needed:

- Document WHY in comments
- Use group vars for groups of hosts with special needs
- Don't scatter overrides across inventory

3. Security Hardening 1 1 HIGH PRIORITY

Critical Issues:

A. Disabled Host Key Checking

Risk: Man-in-the-middle attacks

Fix:

```
[defaults]
host_key_checking = True
# Maintain known_hosts file properly
```

B. Disabled StrictHostKeyChecking

Fix:

- Remove this argument
- Accept host keys on first connection
- Maintain known hosts file

C. PostgreSQL Listening on All Interfaces

```
# group_vars/db_nodes.yml
postgresql_listen_address: "0.0.0.0" # 1 Too permissive
```

Fix:

```
postgresql_listen_address: "127.0.0.1,{{ ansible_default_ipv4.address }}"
# OR use firewall rules to restrict access
```

D. Redis Without Authentication

```
# group_vars/db_nodes.yml
# redis_requirepass: "CHANGEME_STRONG" # ^ Commented out!
```

Fix:

```
# In vault.yml
redis_requirepass: "{{ vault redis password }}"
# In redis role/tasks
- name: Configure Redis authentication
 lineinfile:
    path: /etc/redis/redis.conf
    regexp: '^# requirepass'
   line: "requirepass {{ redis_requirepass }}"
  notify: restart redis
```

E. Hardcoded IP in Variables

```
# group vars/db nodes.yml
redis_additional_bind_ip: "192.168.10.48" # \( \) Violates FQDN policy
```

Fix:

```
redis_additional_bind_ip: "{{ hx hosts ip[hx hosts fqdn['hx-sqldb-server']] }}"
```

4. Missing Role Documentation 1



Issue: Zero or very few role README.md files

Impact:

- New team members struggle to understand roles
- No variable documentation
- No usage examples
- Maintenance difficulties

Recommendation:

Create standardized README.md for each role:

```
# Role: <role name>
## Purpose
Brief description (1-2 sentences)
## Supported OS
- Ubuntu 24.04 (tested)
- Other versions: list or "not tested"
## Requirements
### System Requirements
- Minimum RAM: XGB
- Disk space: XGB
- Network: open ports list
### Dependencies
- Must run after: `role_name`
- Requires: list packages/services
## Variables
### Required Variables
| Variable | Description | Example |
|-----|
| `var_name` | What it does | `value` |
### Optional Variables
| Variable | Description | Default |
|-----|
| `var_name` | What it does | `value` |
## Usage Example
\`\`\`yaml
- hosts: target_group
 roles:
   - role: role name
     var_name: custom_value
/,/,/,
## Tags
- `tag name` - Description of what this tag controls
## Handlers
- `handler_name` - Triggered when: description
## Files & Templates
- `template.j2` - Purpose and usage
## Testing
\`\`\`bash
# How to test this role
ansible-playbook -i inventory/prod.ini playbooks/test-role.yml --check
1,1,1,
## Maintenance Notes
Special considerations for updates/patches
## Author Information
Team name - contact@email.com
```

Priority Roles to Document First:

- 1. base-setup Used by all hosts
- 2. postgresql Critical data layer
- 3. redis Critical cache layer
- 4. orchestrator_* Core application logic
- 5. ollama LLM infrastructure

5. Missing Rollback Procedures 1

Current State:

- Comprehensive recovery for MCP (recover-mcp-server.yml)
- No general rollback playbooks
- No documented rollback procedures

Risk:

- Failed deployments leave system in broken state
- No easy way to revert to previous working version
- Manual recovery required

Recommendation:

Create rollback playbooks:

```
# playbooks/rollback-orchestrator.yml
- name: Rollback orchestrator to previous version
 hosts: orchestrator nodes
 become: yes
  vars_prompt:
    - name: rollback_version
      prompt: "Enter version to rollback to"
      private: no
  pre_tasks:
    - name: Verify rollback version exists
        path: "/opt/hx-citadel-shield/backups/{{ rollback version }}"
      register: backup_check
      failed_when: not backup_check.stat.exists
  tasks:
    - name: Stop current services
      systemd:
        name: "{{ item }}"
        state: stopped
      loop: "{{ orchestrator services }}"
    - name: Restore previous version
      copy:
        src: "/opt/hx-citadel-shield/backups/{{ rollback_version }}/"
        dest: "/opt/hx-citadel-shield/current/"
        remote_src: yes
    - name: Start services
      systemd:
       name: "{{ item }}"
        state: started
      loop: "{{ orchestrator_services }}"
  post_tasks:
    - name: Verify services are healthy
      uri:
        url: "http://localhost:8000/health"
       status code: 200
      retries: 3
      delay: 5
```

Backup Strategy to Implement:

- 1. Pre-deployment backups
- 2. Version tagging
- 3. Backup retention policy
- 4. Backup verification

6. No CI/CD Workflow for Playbooks 1



Current State:

- Tests exist (see TESTS ANALYSIS.md)
- Type checking configured
- No GitHub Actions workflow for playbook testing

- Syntax validation (ansible-playbook --syntax-check)

- Lint checks (ansible-lint)
- Dry run testing (--check mode)
- Integration tests

Recommendation:

Create .github/workflows/ansible-ci.yml :

```
name: Ansible CI
 pull request:
    paths:
     - '**.yml'
      - '**.yaml'
      - 'roles/**'
      - 'playbooks/**'
  push:
    branches:
     - main
jobs:
 lint:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - name: Set up Python
        uses: actions/setup-python@v4
        with:
          python-version: '3.12'
      - name: Install dependencies
        run: |
          pip install ansible ansible-lint yamllint
          ansible-galaxy install -r requirements.yml
      - name: Run yamllint
        run: yamllint .
      - name: Run ansible-lint
        run: ansible-lint
      - name: Syntax check all playbooks
          for playbooks/*.yml; do
            ansible-playbook --syntax-check "$playbook"
          done
      - name: Run check-fqdn.sh
        run: ./scripts/check-fqdn.sh
      - name: Type validation
        run: ./scripts/validate-types.sh
  test:
    runs-on: ubuntu-latest
    needs: lint
    steps:
      - uses: actions/checkout@v3
      - name: Run unit tests
        run: |
          pip install -r requirements-dev.txt
          pytest tests/unit/ -v
      - name: Run integration tests (if safe)
        run:
          # Dry run on test inventory
          ansible-playbook site.yml -i inventory/test.ini --check
```

7. Inventory Environment Separation 1

Current State:

- inventory/prod.ini Production inventory
- inventory/hx-qwui.ini Appears to be a single-host inventory

Missing:

- Development inventory
- Staging inventory
- Testing inventory

Risk:

- Accidental production deployments during testing
- Can't test changes in isolated environment

Recommendation:

Create environment-specific inventories:

```
inventory/
  prod.ini # Production
staging.ini # Staging environment
dev.ini # Development
prod.ini
    test.ini
                        # Testing/CI
    group_vars/
                       # Prod-specific vars
    ├── prod/
         staging/ # Staging-specific vars

dev/ # Dev-specific vars
                        # Dev-specific vars
        dev/
```

Usage:

```
# Development
ansible-playbook site.yml -i inventory/dev.ini
# Staging (with vault)
ansible-playbook site.yml -i inventory/staging.ini --vault-password-file ~/.vault pass
staging
# Production (requires explicit confirmation)
ansible-playbook site.yml -i inventory/prod.ini --vault-password-file ~/.vault pass pr
```

8. No Monitoring/Alerting Configuration 1



Current State:

- monitoring nodes group exists
- hx-metrics-server host defined
- observability role exists
- No visible monitoring playbook in main deployment

Missing:

- Prometheus/Grafana deployment
- Alert definitions
- Service health checks
- Metrics collection configuration

Recommendation:

1. Add monitoring playbook to site.yml:

```
# site.yml
- import_playbook: playbooks/deploy-base.yml
import playbook: playbooks/deploy-monitoring.yml # Add this
- import_playbook: playbooks/deploy-db.yml
# ... rest
```

1. Create comprehensive monitoring:

```
# playbooks/deploy-monitoring.yml
- name: Deploy monitoring stack
 hosts: monitoring nodes
 become: yes
 roles:
   - prometheus
   - grafana
   - alertmanager
    - node_exporter # For all hosts
```

1. Define health checks:

```
# Integrate with smoke-tests.yml
- name: Register services with health checks
 uri:
    url: "http://{{ monitoring_server }}/api/v1/targets"
    method: POST
    body_format: json
    body:
     targets:
       - "{{ inventory_hostname }}:9090" # Prometheus
        - "{{ inventory_hostname }}:5432" # PostgreSQL
        - "{{ inventory hostname }}:6379" # Redis
```

9. Template Version Control 1



Issue: Templates in roles don't have version indicators

Example:

```
roles/base-setup/templates/health-check.sh.j2
roles/base-setup/templates/manage-services.sh.j2
```

Problem:

- Hard to track template changes
- No way to know which version is deployed
- Difficult to debug template-related issues

Recommendation:

Add version headers to templates:

```
#!/bin/bash
# Template: health-check.sh.j2
# Version: 1.2.0
# Last Modified: 2025-10-12
# Description: Health check script for HX-Citadel services
# Changes:
# 1.2.0 - Added circuit breaker status check
# 1.1.0 - Added Redis connectivity test
# Generated by Ansible on {{ ansible_date_time.date }}
# Host: {{ inventory_hostname }}
# Managed: true - DO NOT EDIT MANUALLY
```

10. No Dependency Documentation 1

Issue: Role dependencies not clearly documented

Current: Some roles clearly depend on others, but it's not documented

Example:

- orchestrator postgresql requires base-setup and postgresql
- orchestrator redis requires redis
- Execution order matters but isn't enforced

Recommendation:

Use meta/main.yml in roles:

```
# roles/orchestrator_postgresql/meta/main.yml
dependencies:
  - role: base-setup
  - role: postgresql
    vars:
      postgresql_ensure_running: yes
galaxy_info:
  author: HX-Citadel Team
  description: PostgreSQL integration for orchestrator
 license: MIT
  min_ansible_version: 2.14
  platforms:
    - name: Ubuntu
      versions:
        - jammy
        - noble
```

Critical Recommendations

Priority 1: Security (Immediate Action Required)

Issue	Risk Level	Action
Disabled host key checking	CRITICAL	Enable immediately or docu- ment risk acceptance
Redis without auth	CRITICAL	Enable requirepass via vault
PostgreSQL on 0.0.0.0	HIGH	Restrict to specific interfaces
StrictHostKeyChecking=no	CRITICAL	Remove from inventory
Vault file status unclear	HIGH	Investigate and document

Priority 2: Documentation (High Priority)

Task	Effort	Impact
Create role README files	High	High
Document vault procedures	Low	Critical
Create rollback playbooks	Medium	High
Add inline comments to complex playbooks	Low	Medium

Priority 3: Operational Excellence (Medium Priority)

Task	Effort	Impact
Implement CI/CD for play- books	Medium	High
Create environment-specific inventories	Low	High
Add monitoring deployment	High	High
Version control for templates	Low	Medium
Standardize SSH keys	Low	Medium

Priority 4: Enhancement (Lower Priority)

Task	Effort	Impact
Pin external role versions	Low	Medium
Add pre-task validation	Medium	Medium
Create change log system	Low	Low
Implement backup automation	High	High

Security Considerations

1. Secrets Management

Current State:

- Vault configured: ~/.ansible_vault_pass
- vault.yml exists but status unclear
- Some passwords commented out (Redis)

Best Practices:

A. Vault Structure

```
# group_vars/all/vault.yml (encrypted)
---
# Database Credentials
vault_db_password: "super_secure_password_here"
vault_db_admin_password: "another_secure_password"

# Redis Credentials
vault_redis_password: "redis_secure_password"

# API Keys
vault_linear_api_token: "lin_api_xxxxx"
vault_slack_webhook: "https://hooks.slack.com/services/xxxxx"

# SSH Keys (if needed)
vault_deploy_ssh_key: |
-----BEGIN OPENSSH PRIVATE KEY-----
...
-----END OPENSSH PRIVATE KEY------
```

B. Using Vault Variables

```
# In roles/redis/tasks/main.yml
- name: Configure Redis password
lineinfile:
    path: /etc/redis/redis.conf
    regexp: '^# requirepass'
    line: "requirepass {{ vault_redis_password }}"
    no_log: true # Prevent password from appearing in logs
    notify: restart redis
```

C. Vault Password Management

```
# Development
export ANSIBLE_VAULT_PASSWORD_FILE=~/.ansible_vault_pass_dev

# Production
export ANSIBLE_VAULT_PASSWORD_FILE=~/.ansible_vault_pass_prod

# Or use password prompt
ansible-playbook site.yml --ask-vault-pass
```

2. SSH Key Management

Recommendations:

1. Use ed25519 keys (modern, secure):

```
ssh-keygen -t ed25519 -C "ansible-deploy@hx-citadel"
```

1. Separate keys per environment:

```
~/.ssh/
|— id_ed25519_dev  # Development
|— id_ed25519_staging  # Staging
|— id_ed25519_prod  # Production (restricted access)
```

1. Key rotation schedule:

- Development: Every 6 months
- Staging: Every 3 months
- Production: Every 3 months or after personnel changes

3. Network Security

Current Issues:

```
postgresql_listen_address: "0.0.0.0" # Too permissive
```

Recommendations:

A. Restrict PostgreSQL

```
# group_vars/db_nodes.yml
postgresql_listen_address: "127.0.0.1,{{ ansible_default_ipv4.address }}"

# Or use specific FQDN
postgresql_listen_address: "127.0.0.1,{{ hx_hosts_ip[hx_hosts_fqdn['hx-sqldb-server']] }}"
```

B. Implement Firewall Rules

```
# In base-setup role
- name: Allow PostgreSQL from specific hosts only
ufw:
    rule: allow
    from_ip: "{{ item }}"
    to_port: '5432'
    proto: tcp
loop:
    - "{{ hx_hosts_ip[hx_hosts_fqdn['hx-orchestrator-server']] }}"
    - "{{ hx_hosts_ip[hx_hosts_fqdn['hx-api-server']] }}"
```

C. Redis Security

```
# Configure Redis to bind specific interfaces
redis_bind_interface: "127.0.0.1 {{ ansible_default_ipv4.address }}"

# Enable authentication
redis_requirepass: "{{ vault_redis_password }}"

# Disable dangerous commands
redis_rename_commands:
    - { command: "CONFIG", rename: "CONFIG_{{ 99999 | random }}" }
    - { command: "FLUSHALL", rename: "FLUSHALL_{{ 99999 | random }}" }
```

4. Privilege Escalation

Current State:

- become: yes used appropriately
- Sudo configuration templated
- Unix socket auth for PostgreSQL

Recommendations:

A. Principle of Least Privilege

```
# Don't do this
become: yes # Root for everything

# Do this instead
become: yes
become_user: postgres # Only escalate to needed user
```

B. Sudo Configuration

```
# templates/sudoers.j2
# Allow app user to manage specific services only
{{ app_user }} ALL=(root) NOPASSWD: /bin/systemctl start hx-*
{{ app_user }} ALL=(root) NOPASSWD: /bin/systemctl stop hx-*
{{ app_user }} ALL=(root) NOPASSWD: /bin/systemctl restart hx-*
{{ app_user }} ALL=(root) NOPASSWD: /bin/systemctl status hx-*

# Deny everything else
{{ app_user }} ALL=!/bin/bash, !/bin/sh
```

5. Audit Logging

Recommendation:

Add audit logging to sensitive operations:

Operational Excellence

1. Deployment Workflow

Recommended Process:

```
# 1. Pre-flight checks
ansible-playbook playbooks/preflight-check.yml -i inventory/prod.ini

# 2. Dry run
ansible-playbook site.yml -i inventory/prod.ini --check --diff

# 3. Limited deployment (test one host)
ansible-playbook site.yml -i inventory/prod.ini --limit hx-test-server

# 4. Gradual rollout
ansible-playbook site.yml -i inventory/prod.ini --limit db_nodes
ansible-playbook site.yml -i inventory/prod.ini --limit vector_nodes
# ... etc

# 5. Full deployment
ansible-playbook site.yml -i inventory/prod.ini

# 6. Post-deployment validation
ansible-playbook playbooks/smoke-tests.yml -i inventory/prod.ini

# 7. Notify team
./scripts/slack-notify.sh "Deployment completed successfully"
```

2. Change Management

Create deployment runbook:

```
# Deployment Runbook
## Pre-Deployment
- [ ] Review changes in PR
- [ ] Run tests locally
- [ ] Backup current configuration
- [ ] Notify team in Slack #deployments
- [ ] Schedule maintenance window if needed
## Deployment Steps
1. [ ] Run preflight checks
2. [ ] Execute dry run
3. [ ] Deploy to test environment
4. [ ] Verify test deployment
5. [ ] Deploy to production
6. [ ] Run smoke tests
7. [ ] Monitor for 15 minutes
## Post-Deployment
- [ ] Verify all services healthy
- [ ] Check logs for errors
- [ ] Update documentation
- [ ] Close Linear issue
- [ ] Notify team of completion
## Rollback Procedure
If deployment fails:
1. [ ] Run rollback playbook
2. [ ] Verify rollback successful
3. [ ] Investigate failure
4. [ ] Document in Linear
```

3. Monitoring Integration

Recommended additions to roles:

```
# In each service role
- name: Register service with monitoring
uri:
    url: "http://{{ monitoring_server }}/api/register"
    method: POST
    body_format: json
    body:
        service_name: "{{ service_name }}"
        host: "{{ inventory_hostname }}"
        port: "{{ service_port }}"
        health_check_url: "http://{{ inventory_hostname }}:{{ service_port }}/health"
    delegate_to: localhost
    when: monitoring_enabled | default(true)
```

4. Performance Optimization

Current optimizations:

- V SSH multiplexing
- Pipelining
- V Fact caching

Additional recommendations:

A. Parallel Execution

```
# ansible.cfg
[defaults]
forks = 10  # Increase from default 5 for faster execution
```

B. Selective Fact Gathering

```
# For playbooks that don't need all facts
- name: Quick deployment
hosts: all
gather_facts: no # Skip if not needed
tasks:
    - name: Gather minimal facts
    setup:
    gather_subset:
        - '!all'
        - 'network'
when: network_info_needed
```

C. Async Tasks

```
# For long-running operations
- name: Download large model
get_url:
    url: "{{ model_url }}"
    dest: "/opt/models/"
async: 3600 # 1 hour timeout
poll: 0 # Don't wait
register: download_job

- name: Wait for download
async_status:
    jid: "{{ download_job.ansible_job_id }}"
register: job_result
until: job_result.finished
retries: 360
delay: 10
```

5. Disaster Recovery

Create DR playbooks:

```
# playbooks/disaster-recovery.yml
- name: Full system recovery
 hosts: all
 become: yes
 serial: 1 # One host at a time
 pre_tasks:
    - name: Verify backup exists
      stat:
        path: "/backups/{{ recovery_date }}"
      register: backup_check
      failed_when: not backup_check.stat.exists
  roles:
   - restore system
    - restore database
    - restore_application
  post_tasks:
    - name: Verify services
      systemd:
       name: "{{ item }}"
       state: started
      loop: "{{ critical_services }}"
    - name: Run smoke tests
      import_tasks: ../playbooks/smoke-tests.yml
```

6. Documentation Standards

Maintain these documents:

- 1. **README.md** (Root) Already exists
- 2. **CONTRIBUTING.md** (How to contribute)
- 3. ARCHITECTURE.md (System architecture)

- 4. RUNBOOKS/ (Operational procedures)
 - deployment.md
 - rollback.md
 - disaster-recovery.md
 - troubleshooting.md
- 5. CHANGELOG.md (Version history)
- 6. Role README files (Per role documentation)

Conclusion

Overall Assessment: 8.5/10

The HX-Citadel Ansible infrastructure is **well-architected and production-ready** with strong adherence to best practices. The FQDN management system, modular role design, and comprehensive testing framework are exemplary.

Critical Actions Required

1. Security (Within 1 week):

- V Enable host key checking
- Configure Redis authentication
- Restrict PostgreSQL listening
- V Investigate vault.yml.broken

2. Documentation (Within 2 weeks):

- Create role README files (start with critical roles)
- Document vault recovery procedures
- Create deployment runbooks

3. Operational (Within 1 month):

- Implement CI/CD for playbooks
- Create rollback procedures
- Add monitoring deployment
- Standardize SSH keys

Final Thoughts

This infrastructure demonstrates **strong engineering discipline** and thoughtful design. The identified improvements are primarily about:

- Hardening security (from "good" to "excellent")
- Improving documentation (for team growth)
- Adding operational safeguards (rollback, DR)
- Enforcing consistency (SSH keys, environments)

With the recommended changes, this would easily be a **9.5/10** infrastructure.

Next Steps

- 1. Share this analysis with the team
- 2. Prioritize recommendations based on your context
- 3. Create Linear issues for tracking improvements
- 4. Schedule security review for critical items

5. **Implement incrementally** - don't try to fix everything at once

Appendix: Quick Reference

Common Commands

```
# Syntax check
ansible-playbook playbooks/deploy-base.yml --syntax-check
# Dry run
ansible-playbook site.yml -i inventory/prod.ini --check --diff
# Run specific tags
ansible-playbook site.yml -i inventory/prod.ini --tags "db,vector"
# Limit to specific hosts
ansible-playbook site.yml -i inventory/prod.ini --limit "llm nodes"
# Vault operations
ansible-vault edit group_vars/all/vault.yml
ansible-vault view group_vars/all/vault.yml
ansible-vault rekey group_vars/all/vault.yml
# List hosts
ansible-inventory -i inventory/prod.ini --list
ansible-inventory -i inventory/prod.ini --graph
# Ad-hoc commands
ansible all -i inventory/prod.ini -m ping
ansible db_nodes -i inventory/prod.ini -m shell -a "systemctl status postgresql"
```

Directory Structure Reference

```
hx-citadel-ansible/
ansible.cfg
                         # Main configuration
├── site.yml
                         # Master playbook
  requirements.yml
                       # External dependencies
  inventory/
                         # Host definitions
   ☐ prod.ini
group vars/
                         # Group variables
Ħ

☐ all/
\overline{\mathbb{D}}
   main.yml
   fqdn_map.yml # + FQDN mappings
vault.yml # Encrypted secret
# Encrypted secrets
    db_nodes.yml
\overline{\mathbb{D}}
☐ llm_nodes.yml
                         # Host-specific variables
  host_vars/
                         # Ansible roles (24 custom)
   roles/
                      # Playbooks (26 total)
   playbooks/
                          # Utility scripts (9 scripts)
   scripts/
                          # Test suite
   tests/
   docs/
                           # Documentation
```

Analysis Complete

For questions or clarifications, refer to the project README.md or contact the infrastructure team.